



Machine Learning Model For Insurance Premium Prediction

Project Objective:

Accurately predict health insurance premiums for individuals using machine learning.

Enable insurers to offer fair, personalized pricing while managing risk effectively.

Identify key health and demographic factors that influence insurance costs.

Goals:

Prepare data: handle missing values, outliers.

Feature engineering: BMI, AgeGroup, ChronicConditionCount, interactions, HealthRiskScore.

Scale numeric features, encode categorical/binary features.

Model Selection:

Linear Regression: baseline, interpretable.

Tree-based models: Decision Tree, Random Forest, Gradient Boosting for non-linear relationships.

Evaluation:

Cross-validation (k-fold) for robust performance.

Metrics: RMSE, MAE, R².

Optional: prediction intervals for reliability.

Interpretability:

Feature importance (tree-based or SHAP).

Identify key risk factors affecting premiums.

```
In [58]: # importing libraries
import pandas as pd
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
import xgboost as xgb
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns
import joblib
```

```
In [4]: # loading data
df = pd.read_csv('insurance.csv')
df.head()
```

```
Out[4]:   Age  Diabetes  BloodPressureProblems  AnyTransplants  AnyChronicDiseases
0      45         0                  0                 0                  0
1      60         1                  0                 0                  0
2      36         1                  1                 0                  0
3      52         1                  1                 0                  1
4      38         0                  0                 0                  1
```

```
In [5]: print(df.isnull().sum())
```

```
Age                      0
Diabetes                  0
BloodPressureProblems     0
AnyTransplants             0
AnyChronicDiseases        0
Height                     0
Weight                     0
KnownAllergies             0
HistoryOfCancerInFamily    0
NumberOfMajorSurgeries      0
PremiumPrice                0
dtype: int64
```

```
In [6]: # Check for unusual values
print(df.describe())
```

	Age	Diabetes	BloodPressureProblems	AnyTransplants	\
count	986.000000	986.000000	986.000000	986.000000	
mean	41.745436	0.419878	0.468560	0.055781	
std	13.963371	0.493789	0.499264	0.229615	
min	18.000000	0.000000	0.000000	0.000000	
25%	30.000000	0.000000	0.000000	0.000000	
50%	42.000000	0.000000	0.000000	0.000000	
75%	53.000000	1.000000	1.000000	0.000000	
max	66.000000	1.000000	1.000000	1.000000	
	AnyChronicDiseases	Height	Weight	KnownAllergies	\
count	986.000000	986.000000	986.000000	986.000000	
mean	0.180527	168.182556	76.950304	0.215010	
std	0.384821	10.098155	14.265096	0.411038	
min	0.000000	145.000000	51.000000	0.000000	
25%	0.000000	161.000000	67.000000	0.000000	
50%	0.000000	168.000000	75.000000	0.000000	
75%	0.000000	176.000000	87.000000	0.000000	
max	1.000000	188.000000	132.000000	1.000000	
	HistoryOfCancerInFamily	NumberOfMajorSurgeries	PremiumPrice		
count	986.000000	986.000000	986.000000		
mean	0.117647	0.667343	24336.713996		
std	0.322353	0.749205	6248.184382		
min	0.000000	0.000000	15000.000000		
25%	0.000000	0.000000	21000.000000		
50%	0.000000	1.000000	23000.000000		
75%	0.000000	1.000000	28000.000000		
max	1.000000	3.000000	40000.000000		

Though we do not have missing values, let's add some logic for it

```
In [8]: # Impute missing values
# Numeric features

numeric_features = ['Age', 'Height', 'Weight', 'NumberOfMajorSurgeries']
num_imputer = SimpleImputer(strategy='median') # or strategy='mean'
df[numeric_features] = num_imputer.fit_transform(df[numeric_features])
```

```
In [9]: # Binary features
binary_features = ['Diabetes', 'BloodPressureProblems', 'AnyTransplants',
                   'AnyChronicDiseases', 'KnownAllergies', 'HistoryOfCancerInF

binary_imputer = SimpleImputer(strategy='most_frequent')
df[binary_features] = binary_imputer.fit_transform(df[binary_features])
```

Feature Engineering

```
In [10]: # Create BMI
# Convert height from cm to meters
df['Height_m'] = df['Height'] / 100
```

```

df['BMI'] = df['Weight'] / (df['Height_m'] ** 2)

In [11]: # AgeGroup
bins = [17, 30, 45, 66] # define bins
labels = ['18-30', '31-45', '46-66']
df['AgeGroup'] = pd.cut(df['Age'], bins=bins, labels=labels)

In [12]: # ChronicConditionCount
df['ChronicConditionCount'] = df['Diabetes'] + df['BloodPressureProblems'] + c

```

Scaling and Encoding

```

In [14]: # Scaling numerical features
num_features = ['Age', 'Height', 'Weight', 'BMI', 'NumberOfMajorSurgeries']
scaler = StandardScaler()
df[num_features] = scaler.fit_transform(df[num_features])

In [15]: # Encoding categorical features
df = pd.get_dummies(df, columns=['AgeGroup'], drop_first=True)

In [18]: # Weight-to-Height Ratio
df['Weight_Height_Ratio'] = df['Weight'] / df['Height']

In [19]: # Chronic Conditions × Age
df['Age_ChronicInteraction'] = df['Age'] * df['ChronicConditionCount']

In [20]: # Surgeries × Chronic Conditions
df['Surgery_ChronicInteraction'] = df['NumberOfMajorSurgeries'] * df['ChronicC']

In [21]: # Health Risk Scores
df['HealthRiskScore'] = (2*df['Diabetes'] + 1.5*df['BloodPressureProblems'] +
2.5*df['AnyChronicDiseases'] + df['HistoryOfCancerInFam']

In [22]: # Age Buckets Beyond Simple Groups
df['AgeQuantile'] = pd.qcut(df['Age'], 4, labels=False) # 4 quantiles

In [23]: # Family/Medical History Features
df['FamilyHealthRisk'] = df['AnyChronicDiseases'] + df['HistoryOfCancerInFam']

In [24]: # Lifestyle
df['Obese'] = (df['BMI'] >= 30).astype(int)

In [25]: # Surgery Burden
df['SurgeryBurden'] = pd.cut(df['NumberOfMajorSurgeries'], bins=[-1,0,1,3], la

In [26]: print(df.dtypes)

```

```
Age                      float64
Diabetes                  int64
BloodPressureProblems    int64
AnyTransplants            int64
AnyChronicDiseases        int64
Height                     float64
Weight                     float64
KnownAllergies             int64
HistoryOfCancerInFamily   int64
NumberOfMajorSurgeries     float64
PremiumPrice                int64
Height_m                   float64
BMI                        float64
ChronicConditionCount     int64
AgeGroup_31-45              bool
AgeGroup_46-66              bool
Weight_Height_Ratio         float64
Age_ChronicInteraction     float64
Surgery_ChronicInteraction float64
HealthRiskScore             float64
AgeQuantile                 int64
FamilyHealthRisk            int64
Obese                       int64
SurgeryBurden               category
dtype: object
```

ML Modeling - Prepare Data for Modeling

```
In [27]: X = df.drop('PremiumPrice', axis=1)
y = df['PremiumPrice']
```

```
In [28]: # Handle categorical features
X = pd.get_dummies(X, columns=['SurgeryBurden'], drop_first=True)
```

```
In [30]: # Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Linear Regression (Baseline Model)

```
In [32]: lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
```

```
Out[32]: ▾ LinearRegression ⓘ ?
```

```
LinearRegression()
```

```
In [34]: # Evaluate performance
y_pred = lr_model.predict(X_test)
print("RMSE:", mean_squared_error(y_test, y_pred))
print("MAE:", mean_absolute_error(y_test, y_pred))
print("R2 Score:", r2_score(y_test, y_pred))
```

RMSE: 9264384.065560339
 MAE: 2148.8645406668325
 R2 Score: 0.7827444281107023

```
In [35]: # Check coefficients for feature importance
feature_importance = pd.DataFrame({'Feature': X.columns, 'Coefficient': lr_mod
print(feature_importance.sort_values(by='Coefficient', key=abs, ascending=False))
```

	Feature	Coefficient
3	AnyTransplants	7462.053921
14	AgeGroup_46-66	7101.380932
13	AgeGroup_31-45	5512.071540
0	Age	3342.708470
6	Weight	2023.043965
11	BMI	-1211.228788
20	FamilyHealthRisk	1050.188331
19	AgeQuantile	-804.352492
1	Diabetes	-739.894757
5	Height	-723.841933
22	SurgeryBurden_Low	650.277396
4	AnyChronicDiseases	635.906732
16	Age_ChronicInteraction	-487.444514
8	HistoryOfCancerInFamily	414.281598
9	NumberOfMajorSurgeries	-328.608763
18	HealthRiskScore	326.844304
12	ChronicConditionCount	-235.597766
2	BloodPressureProblems	-131.609741
7	KnownAllergies	-116.848513
10	Height_m	-73.057603
23	SurgeryBurden_High	60.534161
15	Weight_Height_Ratio	-35.654332
17	Surgery_ChronicInteraction	8.729428
21	Obese	0.000000

Decision Tree Regressor

```
In [37]: dt_model = DecisionTreeRegressor(max_depth=5, random_state=42) # max_depth to
dt_model.fit(X_train, y_train)
```

Out[37]:

▾ DecisionTreeRegressor
 ? 

```
DecisionTreeRegressor(max_depth=5, random_state=42)
```

```
In [39]: # Evaluate performance
y_pred_dt = dt_model.predict(X_test)
print("Decision Tree RMSE:", mean_squared_error(y_test, y_pred_dt))
print("Decision Tree R2:", r2_score(y_test, y_pred_dt))
```

```
Decision Tree RMSE: 5689461.284878939
Decision Tree R2: 0.8665785921178095
```

```
In [40]: # Feature importance
```

```
feature_importance_dt = pd.DataFrame({'Feature': X.columns, 'Importance': dt_m
print(feature_importance_dt.sort_values(by='Importance', ascending=False))
```

	Feature	Importance
0	Age	0.704445
3	AnyTransplants	0.113172
6	Weight	0.045837
4	AnyChronicDiseases	0.026759
20	FamilyHealthRisk	0.022815
15	Weight_Height_Ratio	0.021126
9	NumberOfMajorSurgeries	0.018621
2	BloodPressureProblems	0.015430
5	Height	0.013524
11	BMI	0.009436
17	Surgery_ChronicInteraction	0.003427
14	AgeGroup_46-66	0.003241
19	AgeQuantile	0.002165
1	Diabetes	0.000000
7	KnownAllergies	0.000000
10	Height_m	0.000000
12	ChronicConditionCount	0.000000
13	AgeGroup_31-45	0.000000
8	HistoryOfCancerInFamily	0.000000
16	Age_ChronicInteraction	0.000000
18	HealthRiskScore	0.000000
21	Obese	0.000000
22	SurgeryBurden_Low	0.000000
23	SurgeryBurden_High	0.000000

Random Forest Regressor

```
In [42]: rf_model = RandomForestRegressor(n_estimators=200, max_depth=8, random_state=42)
rf_model.fit(X_train, y_train)
```

```
Out[42]:
```

```
RandomForestRegressor
RandomForestRegressor(max_depth=8, n_estimators=200, random_state=42)
```

```
In [44]: # Evaluate performance
```

```
y_pred_rf = rf_model.predict(X_test)
print("Random Forest RMSE:", mean_squared_error(y_test, y_pred_rf))
print("Random Forest R2:", r2_score(y_test, y_pred_rf))
```

```
Random Forest RMSE: 5699235.970613663
Random Forest R2: 0.8663493696541996
```

```
In [45]: # Feature importance
```

```
feature_importance_rf = pd.DataFrame({'Feature': X.columns, 'Importance': rf_m
```

```
print(feature_importance_rf.sort_values(by='Importance', ascending=False))

      Feature  Importance
0            Age     0.608920
3      AnyTransplants    0.091370
6          Weight     0.060870
15   Weight_Height_Ratio    0.052651
11            BMI     0.027376
9  NumberOfMajorSurgeries    0.021621
4      AnyChronicDiseases    0.019715
20    FamilyHealthRisk     0.019657
16  Age_ChronicInteraction    0.018177
17  Surgery_ChronicInteraction    0.013784
10        Height_m     0.013138
5            Height     0.011343
8  HistoryOfCancerInFamily    0.010074
14    AgeGroup_46-66     0.008090
18    HealthRiskScore     0.006473
2    BloodPressureProblems    0.004765
1        Diabetes     0.003183
12  ChronicConditionCount    0.002460
22  SurgeryBurden_Low     0.002109
23  SurgeryBurden_High     0.002024
19        AgeQuantile    0.000858
7      KnownAllergies     0.000830
13    AgeGroup_31-45     0.000511
21         Obese     0.000000
```

Gradient Boosting Machines (XGBoost / LightGBM)

In [47]: `xgb_model = xgb.XGBRegressor(n_estimators=300, learning_rate=0.05, max_depth=5)`

Out[47]:

```
XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, device=None, early_stopping_rounds=None,
             enable_categorical=False, eval_metric=None, feature_type=None,
             feature_weights=None, gamma=None, grow_policy=None,
             importance_type=None, interaction_constraints=None,
             learning_rate=0.05, max_bin=None, max_cat_threshold=None,
             max_delta_step=None, max_depth=5, min_child_weight=1,
             missing='NaN', n_estimators=300, n_jobs=1, objective='reg:squarederror',
             random_state=None, reg_alpha=0.0, reg_lambda=1.0,
             scale_pos_weight=1.0, tree_method='auto')
```

In [49]: `# Evaluate performance`
`y_pred_xgb = xgb_model.predict(X_test)`
`print("XGBoost RMSE:", mean_squared_error(y_test, y_pred_xgb))`
`print("XGBoost R2:", r2_score(y_test, y_pred_xgb))`

```
XGBoost RMSE: 5738102.5  
XGBoost R2: 0.865437924861908
```

```
In [50]: # Feature importance  
xgb_importance = pd.DataFrame({'Feature': X.columns, 'Importance': xgb_model.feature_importances_})  
print(xgb_importance.sort_values(by='Importance', ascending=False))
```

	Feature	Importance
3	AnyTransplants	0.295866
0	Age	0.210831
4	AnyChronicDiseases	0.095494
20	FamilyHealthRisk	0.089955
9	NumberOfMajorSurgeries	0.068664
6	Weight	0.045107
8	HistoryOfCancerInFamily	0.042909
2	BloodPressureProblems	0.023709
15	Weight_Height_Ratio	0.022422
7	KnownAllergies	0.015194
17	Surgery_ChronicInteraction	0.014567
11	BMI	0.014464
5	Height	0.013040
22	SurgeryBurden_Low	0.012340
16	Age_ChronicInteraction	0.010652
1	Diabetes	0.008899
12	ChronicConditionCount	0.006213
18	HealthRiskScore	0.005384
13	AgeGroup_31-45	0.004292
10	Height_m	0.000000
14	AgeGroup_46-66	0.000000
19	AgeQuantile	0.000000
21	Obese	0.000000
23	SurgeryBurden_High	0.000000

Model Comparison

```
In [51]: # Creating the models  
  
models = {  
    'Linear Regression': LinearRegression(),  
    'Decision Tree': DecisionTreeRegressor(max_depth=5, random_state=42),  
    'Random Forest': RandomForestRegressor(n_estimators=200, max_depth=8, random_state=42),  
    'XGBoost': xgb.XGBRegressor(n_estimators=300, learning_rate=0.05, max_depth=5)  
}
```

```
In [53]: results = []  
  
for name, model in models.items():  
    # Train  
    model.fit(X_train, y_train)  
  
    # Predict  
    y_pred = model.predict(X_test)
```

```

# Metrics
rmse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

results.append({
    'Model': name,
    'RMSE': rmse,
    'MAE': mae,
    'R2': r2
})

```

In [54]:

```
# Create comparison table
results_df = pd.DataFrame(results).sort_values(by='RMSE')
print(results_df)
```



	Model	RMSE	MAE	R2
1	Decision Tree	5.689461e+06	1244.826604	0.866579
2	Random Forest	5.699236e+06	1178.751778	0.866349
3	XGBoost	5.738102e+06	1329.295410	0.865438
0	Linear Regression	9.264384e+06	2148.864541	0.782744


In [57]:

```
# Feature Importance

for name, model in models.items():
    if name != 'Linear Regression':
        if hasattr(model, 'feature_importances_'):
            fi = pd.DataFrame({'Feature': X.columns, 'Importance': model.feature_importances_})
            fi = fi.sort_values(by='Importance', ascending=False).head(10)

            plt.figure(figsize=(8,5))
            sns.barplot(x='Importance', y='Feature', data=fi)
            plt.title(f'Top Features - {name}')
            plt.show()
```





Interpretation

Decision Tree, Random Forest, and XGBoost all perform similarly well, with R^2 around 0.865–0.867, meaning they explain about 87% of the variance in the target variable.

Their MAE is much lower (~1,178–1,329) compared to Linear Regression (~2,149), indicating more accurate average predictions.

RMSE is also lower (~5.69–5.74M) for tree-based models than Linear Regression (~9.26M), suggesting fewer extreme errors or outliers in predictions.

Linear Regression underperforms: it has the lowest R^2 (0.78) and the highest errors, both MAE and RMSE, indicating it struggles to capture the complexity in the data.

Conclusion:

Tree-based models (Decision Tree, Random Forest, XGBoost) are clearly superior to Linear Regression for this dataset, with Random Forest having the lowest MAE and a strong balance between accuracy and variance explanation. Linear Regression is not suitable here due to high errors and lower predictive power.

Save the Model

```
In [59]: joblib.dump(rf_model, 'random_forest_insurance_model.pkl')
joblib.dump(dt_model, 'decision_tree_model.pkl')
joblib.dump(xgb_model, 'xgboost_model.pkl')
```

```
joblib.dump(lr_model, 'linear_regression_model.pkl')  
print("Models saved successfully!")
```

Models saved successfully!

Testing the saved model

```
In [60]: loaded_rf_model = joblib.load('random_forest_insurance_model.pkl')  
loaded_dt_model = joblib.load('decision_tree_model.pkl')  
loaded_xgb_model = joblib.load('xgboost_model.pkl')  
loaded_lr_model = joblib.load('linear_regression_model.pkl')
```

```
In [66]: # picking the first record to evaluate the model  
X_sample = X_test.iloc[[0]]  
y_actual = y_test.iloc[0].item()
```

```
In [67]: # Store models in a dictionary for easy iteration  
models = {  
    "Random Forest": loaded_rf_model,  
    "Decision Tree": loaded_dt_model,  
    "XGBoost": loaded_xgb_model,  
    "Linear Regression": loaded_lr_model  
}  
  
# Predict and print results  
for name, model in models.items():  
    y_pred = model.predict(X_sample)  
    print(f"{name} Prediction: {y_pred[0]:.2f}, Actual: {y_actual}")
```

Random Forest Prediction: 29836.17, Actual: 31000
Decision Tree Prediction: 28601.27, Actual: 31000
XGBoost Prediction: 30642.93, Actual: 31000
Linear Regression Prediction: 29980.79, Actual: 31000

```
In [73]: X_sample
```

```
Out[73]:      Age  Diabetes  BloodPressureProblems  AnyTransplants  AnyChronicDi:  
613  0.591458          0                  0                  0
```

1 rows × 24 columns