



```
In [58]: # importing libraries
import pandas as pd
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
import xgboost as xgb
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns
import joblib
```

```
In [4]: # loading data
df = pd.read_csv('insurance.csv')
df.head()
```

```
Out[4]:   Age  Diabetes  BloodPressureProblems  AnyTransplants  AnyChronicDiseases
0      45         0                  0                 0                  0
1      60         1                  0                 0                  0
2      36         1                  1                 0                  0
3      52         1                  1                 0                  1
4      38         0                  0                 0                  1
```

```
In [5]: print(df.isnull().sum())
```

```
Age                      0
Diabetes                  0
BloodPressureProblems     0
AnyTransplants             0
AnyChronicDiseases        0
Height                     0
Weight                     0
KnownAllergies             0
HistoryOfCancerInFamily    0
NumberOfMajorSurgeries      0
PremiumPrice                0
dtype: int64
```

```
In [6]: # Check for unusual values
print(df.describe())
```

	Age	Diabetes	BloodPressureProblems	AnyTransplants	\
count	986.000000	986.000000	986.000000	986.000000	
mean	41.745436	0.419878	0.468560	0.055781	
std	13.963371	0.493789	0.499264	0.229615	
min	18.000000	0.000000	0.000000	0.000000	
25%	30.000000	0.000000	0.000000	0.000000	
50%	42.000000	0.000000	0.000000	0.000000	
75%	53.000000	1.000000	1.000000	0.000000	
max	66.000000	1.000000	1.000000	1.000000	
	AnyChronicDiseases	Height	Weight	KnownAllergies	\
count	986.000000	986.000000	986.000000	986.000000	
mean	0.180527	168.182556	76.950304	0.215010	
std	0.384821	10.098155	14.265096	0.411038	
min	0.000000	145.000000	51.000000	0.000000	
25%	0.000000	161.000000	67.000000	0.000000	
50%	0.000000	168.000000	75.000000	0.000000	
75%	0.000000	176.000000	87.000000	0.000000	
max	1.000000	188.000000	132.000000	1.000000	
	HistoryOfCancerInFamily	NumberOfMajorSurgeries	PremiumPrice		
count	986.000000	986.000000	986.000000		
mean	0.117647	0.667343	24336.713996		
std	0.322353	0.749205	6248.184382		
min	0.000000	0.000000	15000.000000		
25%	0.000000	0.000000	21000.000000		
50%	0.000000	1.000000	23000.000000		
75%	0.000000	1.000000	28000.000000		
max	1.000000	3.000000	40000.000000		

Though we do not have missing values, let's add some logic for it

```
In [8]: # Impute missing values
# Numeric features

numeric_features = ['Age', 'Height', 'Weight', 'NumberOfMajorSurgeries']
num_imputer = SimpleImputer(strategy='median') # or strategy='mean'
df[numeric_features] = num_imputer.fit_transform(df[numeric_features])
```

```
In [9]: # Binary features
binary_features = ['Diabetes', 'BloodPressureProblems', 'AnyTransplants',
                   'AnyChronicDiseases', 'KnownAllergies', 'HistoryOfCancerInF

binary_imputer = SimpleImputer(strategy='most_frequent')
df[binary_features] = binary_imputer.fit_transform(df[binary_features])
```

Feature Engineering

```
In [10]: # Create BMI
# Convert height from cm to meters
df['Height_m'] = df['Height'] / 100
```

```

df['BMI'] = df['Weight'] / (df['Height_m'] ** 2)

In [11]: # AgeGroup
bins = [17, 30, 45, 66] # define bins
labels = ['18-30', '31-45', '46-66']
df['AgeGroup'] = pd.cut(df['Age'], bins=bins, labels=labels)

In [12]: # ChronicConditionCount
df['ChronicConditionCount'] = df['Diabetes'] + df['BloodPressureProblems'] + c

```

Scaling and Encoding

```

In [14]: # Scaling numerical features
num_features = ['Age', 'Height', 'Weight', 'BMI', 'NumberOfMajorSurgeries']
scaler = StandardScaler()
df[num_features] = scaler.fit_transform(df[num_features])

In [15]: # Encoding categorical features
df = pd.get_dummies(df, columns=['AgeGroup'], drop_first=True)

In [18]: # Weight-to-Height Ratio
df['Weight_Height_Ratio'] = df['Weight'] / df['Height']

In [19]: # Chronic Conditions × Age
df['Age_ChronicInteraction'] = df['Age'] * df['ChronicConditionCount']

In [20]: # Surgeries × Chronic Conditions
df['Surgery_ChronicInteraction'] = df['NumberOfMajorSurgeries'] * df['ChronicC']

In [21]: # Health Risk Scores
df['HealthRiskScore'] = (2*df['Diabetes'] + 1.5*df['BloodPressureProblems'] +
2.5*df['AnyChronicDiseases'] + df['HistoryOfCancerInFam']

In [22]: # Age Buckets Beyond Simple Groups
df['AgeQuantile'] = pd.qcut(df['Age'], 4, labels=False) # 4 quantiles

In [23]: # Family/Medical History Features
df['FamilyHealthRisk'] = df['AnyChronicDiseases'] + df['HistoryOfCancerInFam']

In [24]: # Lifestyle
df['Obese'] = (df['BMI'] >= 30).astype(int)

In [25]: # Surgery Burden
df['SurgeryBurden'] = pd.cut(df['NumberOfMajorSurgeries'], bins=[-1,0,1,3], la

In [26]: print(df.dtypes)

```

```
Age                      float64
Diabetes                  int64
BloodPressureProblems    int64
AnyTransplants            int64
AnyChronicDiseases        int64
Height                     float64
Weight                     float64
KnownAllergies             int64
HistoryOfCancerInFamily   int64
NumberOfMajorSurgeries     float64
PremiumPrice                int64
Height_m                   float64
BMI                        float64
ChronicConditionCount     int64
AgeGroup_31-45              bool
AgeGroup_46-66              bool
Weight_Height_Ratio         float64
Age_ChronicInteraction     float64
Surgery_ChronicInteraction float64
HealthRiskScore             float64
AgeQuantile                 int64
FamilyHealthRisk            int64
Obese                       int64
SurgeryBurden               category
dtype: object
```

ML Modeling - Prepare Data for Modeling

```
In [27]: X = df.drop('PremiumPrice', axis=1)
y = df['PremiumPrice']
```

```
In [28]: # Handle categorical features
X = pd.get_dummies(X, columns=['SurgeryBurden'], drop_first=True)
```

```
In [30]: # Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Linear Regression (Baseline Model)

```
In [32]: lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
```

```
Out[32]: ▾ LinearRegression ⓘ ?
```

```
LinearRegression()
```

```
In [34]: # Evaluate performance
y_pred = lr_model.predict(X_test)
print("RMSE:", mean_squared_error(y_test, y_pred))
print("MAE:", mean_absolute_error(y_test, y_pred))
print("R2 Score:", r2_score(y_test, y_pred))
```

RMSE: 9264384.065560339
 MAE: 2148.8645406668325
 R2 Score: 0.7827444281107023

```
In [35]: # Check coefficients for feature importance
feature_importance = pd.DataFrame({'Feature': X.columns, 'Coefficient': lr_mod
print(feature_importance.sort_values(by='Coefficient', key=abs, ascending=False))
```

	Feature	Coefficient
3	AnyTransplants	7462.053921
14	AgeGroup_46-66	7101.380932
13	AgeGroup_31-45	5512.071540
0	Age	3342.708470
6	Weight	2023.043965
11	BMI	-1211.228788
20	FamilyHealthRisk	1050.188331
19	AgeQuantile	-804.352492
1	Diabetes	-739.894757
5	Height	-723.841933
22	SurgeryBurden_Low	650.277396
4	AnyChronicDiseases	635.906732
16	Age_ChronicInteraction	-487.444514
8	HistoryOfCancerInFamily	414.281598
9	NumberOfMajorSurgeries	-328.608763
18	HealthRiskScore	326.844304
12	ChronicConditionCount	-235.597766
2	BloodPressureProblems	-131.609741
7	KnownAllergies	-116.848513
10	Height_m	-73.057603
23	SurgeryBurden_High	60.534161
15	Weight_Height_Ratio	-35.654332
17	Surgery_ChronicInteraction	8.729428
21	Obese	0.000000

Decision Tree Regressor

```
In [37]: dt_model = DecisionTreeRegressor(max_depth=5, random_state=42) # max_depth to control overfitting
dt_model.fit(X_train, y_train)
```

Out[37]: ▾ DecisionTreeRegressor

```
DecisionTreeRegressor(max_depth=5, random_state=42)
```

```
In [39]: # Evaluate performance
y_pred_dt = dt_model.predict(X_test)
print("Decision Tree RMSE:", mean_squared_error(y_test, y_pred_dt))
print("Decision Tree R2:", r2_score(y_test, y_pred_dt))
```

```
Decision Tree RMSE: 5689461.284878939
Decision Tree R2: 0.8665785921178095
```

```
In [40]: # Feature importance
```

```
feature_importance_dt = pd.DataFrame({'Feature': X.columns, 'Importance': dt_m
print(feature_importance_dt.sort_values(by='Importance', ascending=False))
```

	Feature	Importance
0	Age	0.704445
3	AnyTransplants	0.113172
6	Weight	0.045837
4	AnyChronicDiseases	0.026759
20	FamilyHealthRisk	0.022815
15	Weight_Height_Ratio	0.021126
9	NumberOfMajorSurgeries	0.018621
2	BloodPressureProblems	0.015430
5	Height	0.013524
11	BMI	0.009436
17	Surgery_ChronicInteraction	0.003427
14	AgeGroup_46-66	0.003241
19	AgeQuantile	0.002165
1	Diabetes	0.000000
7	KnownAllergies	0.000000
10	Height_m	0.000000
12	ChronicConditionCount	0.000000
13	AgeGroup_31-45	0.000000
8	HistoryOfCancerInFamily	0.000000
16	Age_ChronicInteraction	0.000000
18	HealthRiskScore	0.000000
21	Obese	0.000000
22	SurgeryBurden_Low	0.000000
23	SurgeryBurden_High	0.000000

Random Forest Regressor

```
In [42]: rf_model = RandomForestRegressor(n_estimators=200, max_depth=8, random_state=42)
rf_model.fit(X_train, y_train)
```

```
Out[42]:
```

```
RandomForestRegressor
RandomForestRegressor(max_depth=8, n_estimators=200, random_state=42)
```

```
In [44]: # Evaluate performance
```

```
y_pred_rf = rf_model.predict(X_test)
print("Random Forest RMSE:", mean_squared_error(y_test, y_pred_rf))
print("Random Forest R2:", r2_score(y_test, y_pred_rf))
```

```
Random Forest RMSE: 5699235.970613663
Random Forest R2: 0.8663493696541996
```

```
In [45]: # Feature importance
```

```
feature_importance_rf = pd.DataFrame({'Feature': X.columns, 'Importance': rf_m
```

```
print(feature_importance_rf.sort_values(by='Importance', ascending=False))

      Feature  Importance
0            Age     0.608920
3      AnyTransplants    0.091370
6          Weight     0.060870
15   Weight_Height_Ratio    0.052651
11            BMI     0.027376
9  NumberOfMajorSurgeries    0.021621
4      AnyChronicDiseases    0.019715
20    FamilyHealthRisk     0.019657
16  Age_ChronicInteraction    0.018177
17  Surgery_ChronicInteraction    0.013784
10        Height_m     0.013138
5            Height     0.011343
8  HistoryOfCancerInFamily    0.010074
14    AgeGroup_46-66     0.008090
18    HealthRiskScore     0.006473
2    BloodPressureProblems    0.004765
1        Diabetes     0.003183
12  ChronicConditionCount    0.002460
22    SurgeryBurden_Low     0.002109
23    SurgeryBurden_High     0.002024
19        AgeQuantile     0.000858
7      KnownAllergies     0.000830
13    AgeGroup_31-45     0.000511
21         Obese     0.000000
```

Gradient Boosting Machines (XGBoost / LightGBM)

In [47]: `xgb_model = xgb.XGBRegressor(n_estimators=300, learning_rate=0.05, max_depth=5)`

Out[47]:

```
XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, device=None, early_stopping_rounds=None,
             enable_categorical=False, eval_metric=None, feature_type=None,
             feature_weights=None, gamma=None, grow_policy=None,
             importance_type=None, interaction_constraints=None,
             learning_rate=0.05, max_bin=None, max_cat_threshold=None,
             max_delta_step=None, max_depth=5, min_child_weight=1,
             missing='NaN', n_estimators=300, n_jobs=1, objective='reg:squarederror',
             random_state=None, reg_alpha=0.0, reg_lambda=1.0,
             scale_pos_weight=None, subsample=1.0, tree_method='auto')
```

In [49]: `# Evaluate performance`
`y_pred_xgb = xgb_model.predict(X_test)`
`print("XGBoost RMSE:", mean_squared_error(y_test, y_pred_xgb))`
`print("XGBoost R2:", r2_score(y_test, y_pred_xgb))`

```
XGBoost RMSE: 5738102.5  
XGBoost R2: 0.865437924861908
```

```
In [50]: # Feature importance  
xgb_importance = pd.DataFrame({'Feature': X.columns, 'Importance': xgb_model.feature_importances_})  
print(xgb_importance.sort_values(by='Importance', ascending=False))
```

	Feature	Importance
3	AnyTransplants	0.295866
0	Age	0.210831
4	AnyChronicDiseases	0.095494
20	FamilyHealthRisk	0.089955
9	NumberOfMajorSurgeries	0.068664
6	Weight	0.045107
8	HistoryOfCancerInFamily	0.042909
2	BloodPressureProblems	0.023709
15	Weight_Height_Ratio	0.022422
7	KnownAllergies	0.015194
17	Surgery_ChronicInteraction	0.014567
11	BMI	0.014464
5	Height	0.013040
22	SurgeryBurden_Low	0.012340
16	Age_ChronicInteraction	0.010652
1	Diabetes	0.008899
12	ChronicConditionCount	0.006213
18	HealthRiskScore	0.005384
13	AgeGroup_31-45	0.004292
10	Height_m	0.000000
14	AgeGroup_46-66	0.000000
19	AgeQuantile	0.000000
21	Obese	0.000000
23	SurgeryBurden_High	0.000000

Model Comparison

```
In [51]: # Creating the models  
  
models = {  
    'Linear Regression': LinearRegression(),  
    'Decision Tree': DecisionTreeRegressor(max_depth=5, random_state=42),  
    'Random Forest': RandomForestRegressor(n_estimators=200, max_depth=8, random_state=42),  
    'XGBoost': xgb.XGBRegressor(n_estimators=300, learning_rate=0.05, max_depth=5)  
}
```

```
In [53]: results = []  
  
for name, model in models.items():  
    # Train  
    model.fit(X_train, y_train)  
  
    # Predict  
    y_pred = model.predict(X_test)
```

```

# Metrics
rmse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

results.append({
    'Model': name,
    'RMSE': rmse,
    'MAE': mae,
    'R2': r2
})

```

In [54]:

```
# Create comparison table
results_df = pd.DataFrame(results).sort_values(by='RMSE')
print(results_df)
```

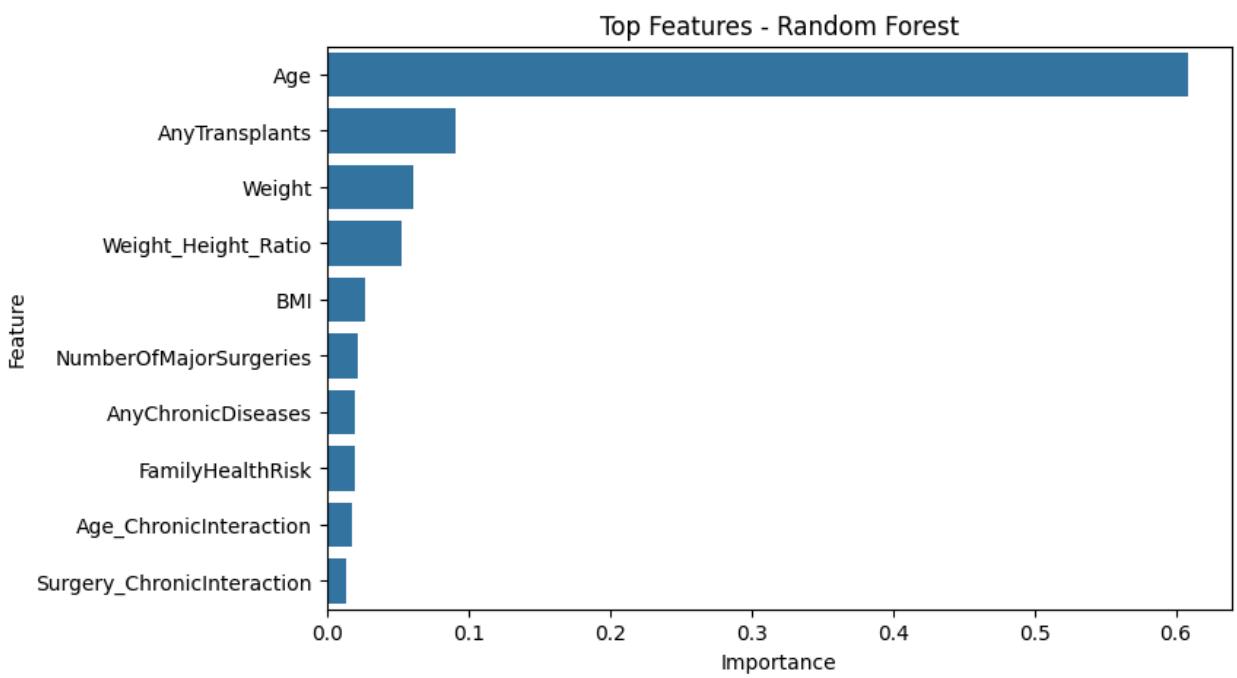
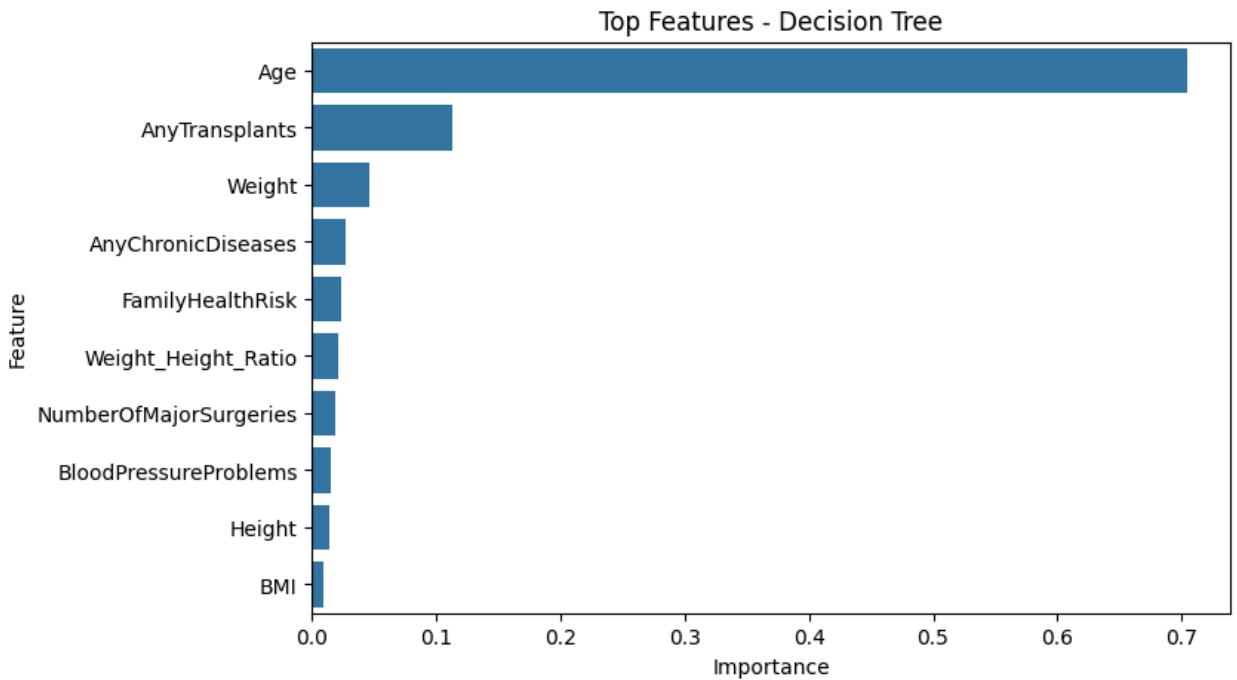
	Model	RMSE	MAE	R2
1	Decision Tree	5.689461e+06	1244.826604	0.866579
2	Random Forest	5.699236e+06	1178.751778	0.866349
3	XGBoost	5.738102e+06	1329.295410	0.865438
0	Linear Regression	9.264384e+06	2148.864541	0.782744

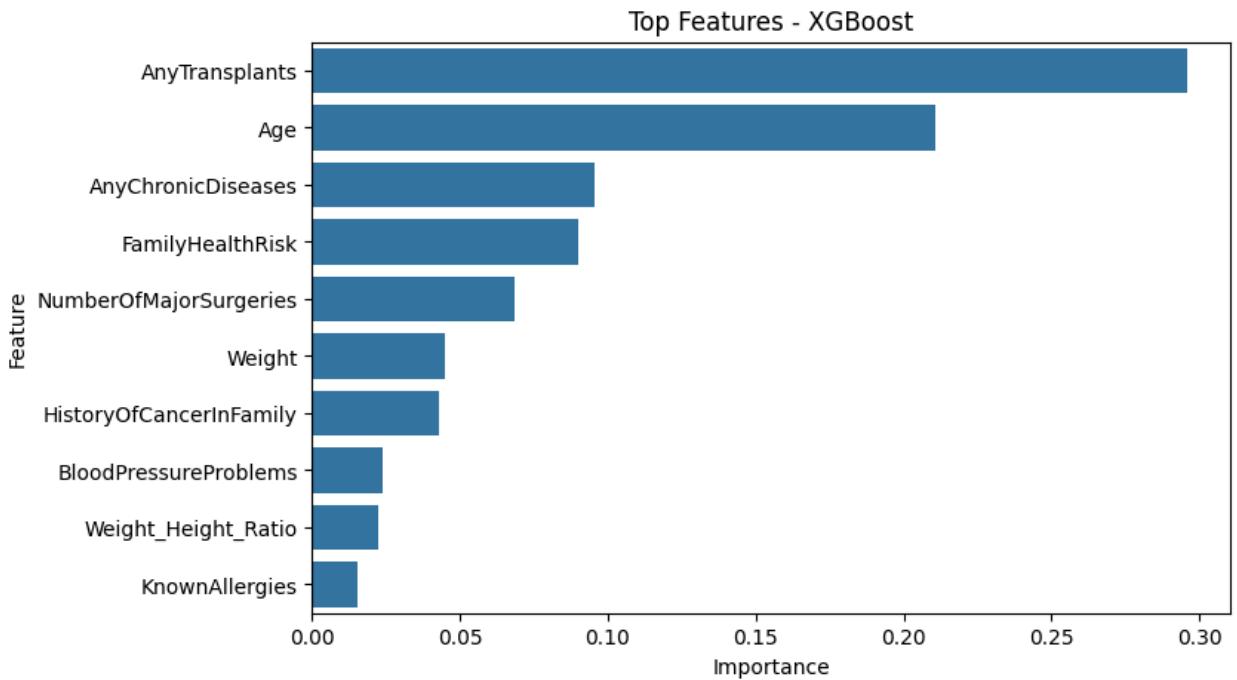
In [57]:

```
# Feature Importance

for name, model in models.items():
    if name != 'Linear Regression':
        if hasattr(model, 'feature_importances_'):
            fi = pd.DataFrame({'Feature': X.columns, 'Importance': model.feature_importances_})
            fi = fi.sort_values(by='Importance', ascending=False).head(10)

            plt.figure(figsize=(8,5))
            sns.barplot(x='Importance', y='Feature', data=fi)
            plt.title(f'Top Features - {name}')
            plt.show()
```





Interpretation

Decision Tree, Random Forest, and XGBoost all perform similarly well, with R^2 around 0.865–0.867, meaning they explain about 87% of the variance in the target variable.

Their MAE is much lower (~1,178–1,329) compared to Linear Regression (~2,149), indicating more accurate average predictions.

RMSE is also lower (~5.69–5.74M) for tree-based models than Linear Regression (~9.26M), suggesting fewer extreme errors or outliers in predictions.

Linear Regression underperforms: it has the lowest R^2 (0.78) and the highest errors, both MAE and RMSE, indicating it struggles to capture the complexity in the data.

Conclusion:

Tree-based models (Decision Tree, Random Forest, XGBoost) are clearly superior to Linear Regression for this dataset, with Random Forest having the lowest MAE and a strong balance between accuracy and variance explanation. Linear Regression is not suitable here due to high errors and lower predictive power.

Save the Model

```
In [59]: joblib.dump(rf_model, 'random_forest_insurance_model.pkl')
joblib.dump(dt_model, 'decision_tree_model.pkl')
joblib.dump(xgb_model, 'xgboost_model.pkl')
```

```
joblib.dump(lr_model, 'linear_regression_model.pkl')  
print("Models saved successfully!")
```

Models saved successfully!

Testing the saved model

```
In [60]: loaded_rf_model = joblib.load('random_forest_insurance_model.pkl')  
loaded_dt_model = joblib.load('decision_tree_model.pkl')  
loaded_xgb_model = joblib.load('xgboost_model.pkl')  
loaded_lr_model = joblib.load('linear_regression_model.pkl')
```

```
In [66]: # picking the first record to evaluate the model  
X_sample = X_test.iloc[[0]]  
y_actual = y_test.iloc[0].item()
```

```
In [67]: # Store models in a dictionary for easy iteration  
models = {  
    "Random Forest": loaded_rf_model,  
    "Decision Tree": loaded_dt_model,  
    "XGBoost": loaded_xgb_model,  
    "Linear Regression": loaded_lr_model  
}  
  
# Predict and print results  
for name, model in models.items():  
    y_pred = model.predict(X_sample)  
    print(f"{name} Prediction: {y_pred[0]:.2f}, Actual: {y_actual}")
```

```
Random Forest Prediction: 29836.17, Actual: 31000  
Decision Tree Prediction: 28601.27, Actual: 31000  
XGBoost Prediction: 30642.93, Actual: 31000  
Linear Regression Prediction: 29980.79, Actual: 31000
```